

# HDFS

---

Senthil Kumar A



# Basics — *optional*

- Program
  - sequence of instructions written to perform a specified task with a computer
  - or a piece of code
- Process
  - an instance of a computer program that is being executed.
  - or a execution of a program
- Daemon Process
  - process which runs in background and has no controlling terminal.
- JVM – Java Virtual Machine
  - program which executes certain programs, namely those containing Java bytecode instructions

# Basics — *optional*

- Client-server Concept

- Client sends requests to one or more servers which in turn accepts, processes them and return the requested information to the client.
- A server might run a software which listens on particular ip and port number for requests
- Examples:
  - Server - web server
  - Client – web browser

# Introduction

- A distributed File System - **STORAGE**

- A File System on multiple machines which sits on native filesystem
  - ext4,ext3

- Hardware Failure

- Due to usage of Commodity machines, failure is a common phenomenon
- Designed for failure

- Large Data Sets

- Small Files Problem Due to NameNode

- Simple Coherency Model

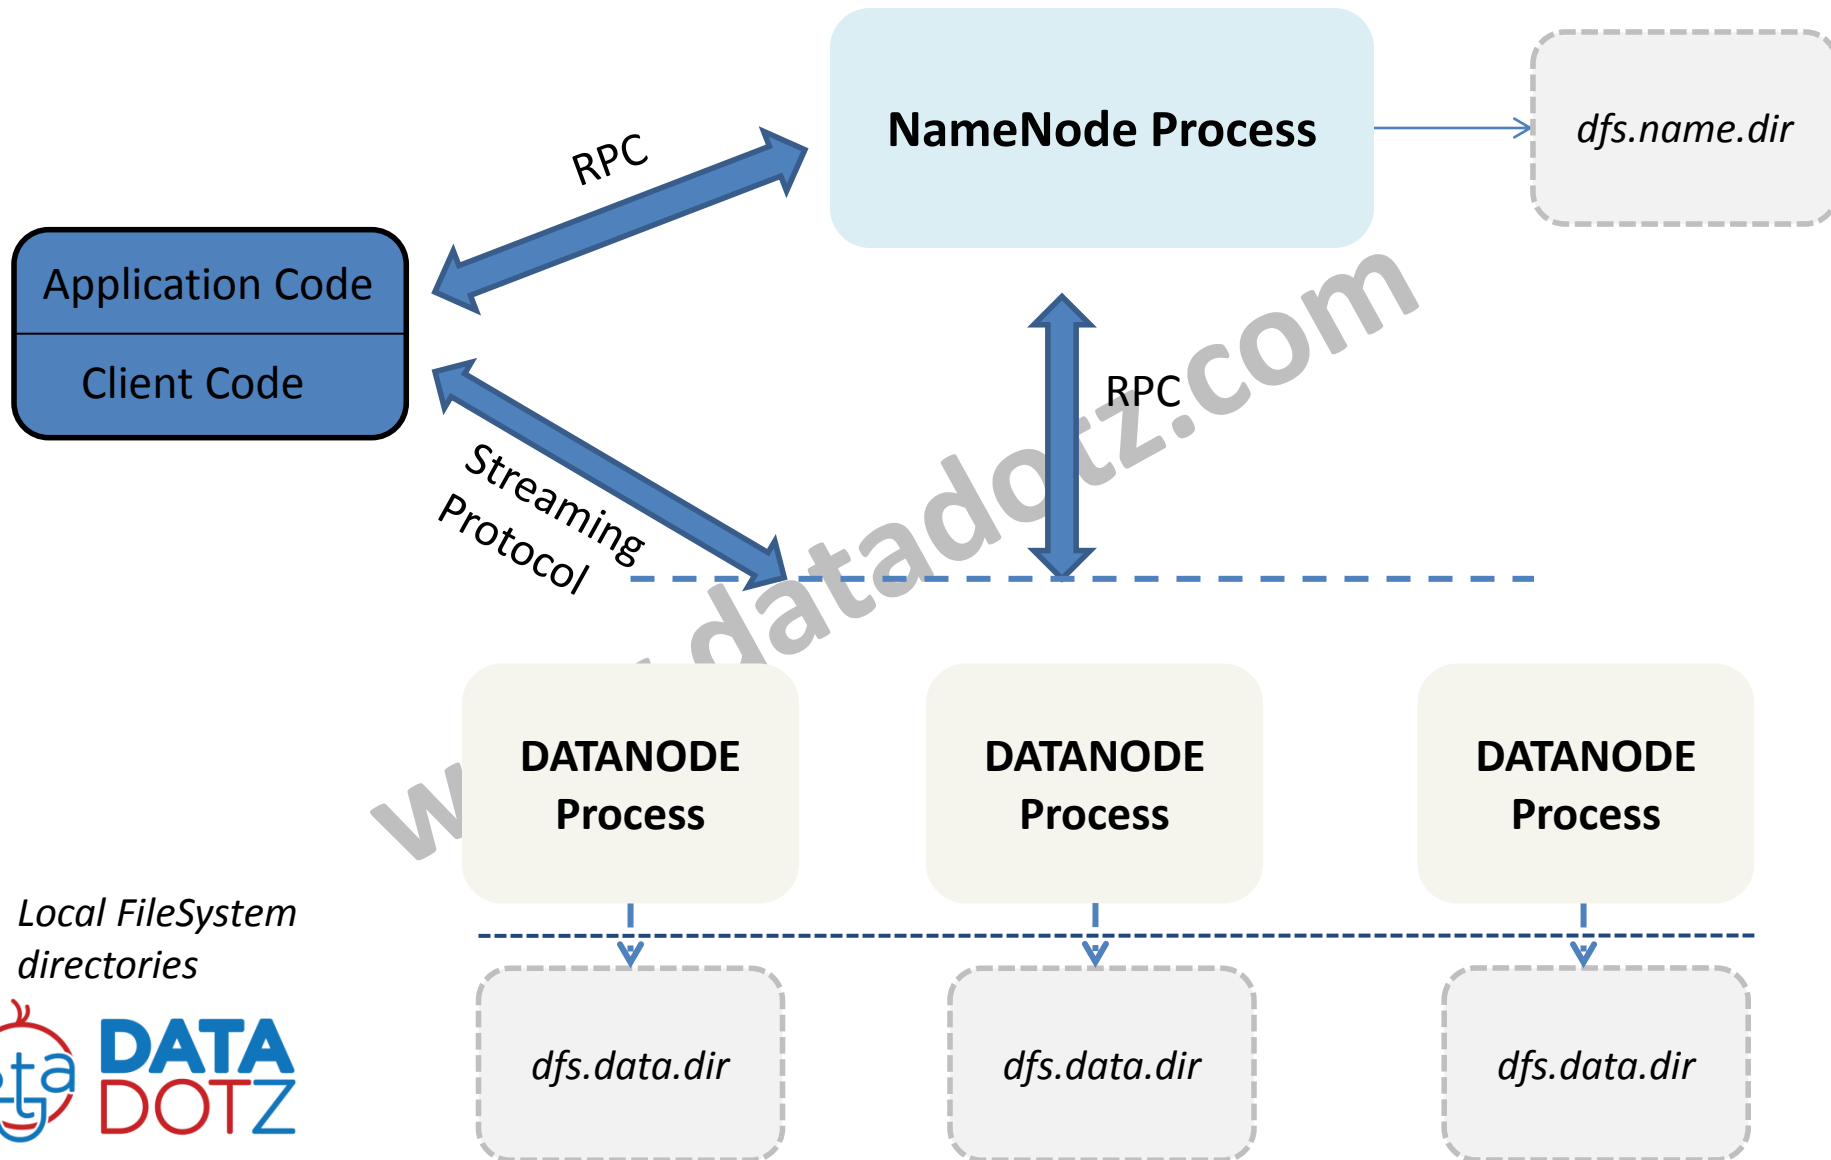
- Write Once , Read Many Times

- Streaming Data Access

- High Throughput instead of low latency access

- ext3? ext4?

# Continued...



# Daemons in Hadoop Core

- NameNode
  - DataNode
  - Secondary NameNode\*
- } **HDFS**
- JobTracker\*
  - TaskTracker\*
- } **MR**

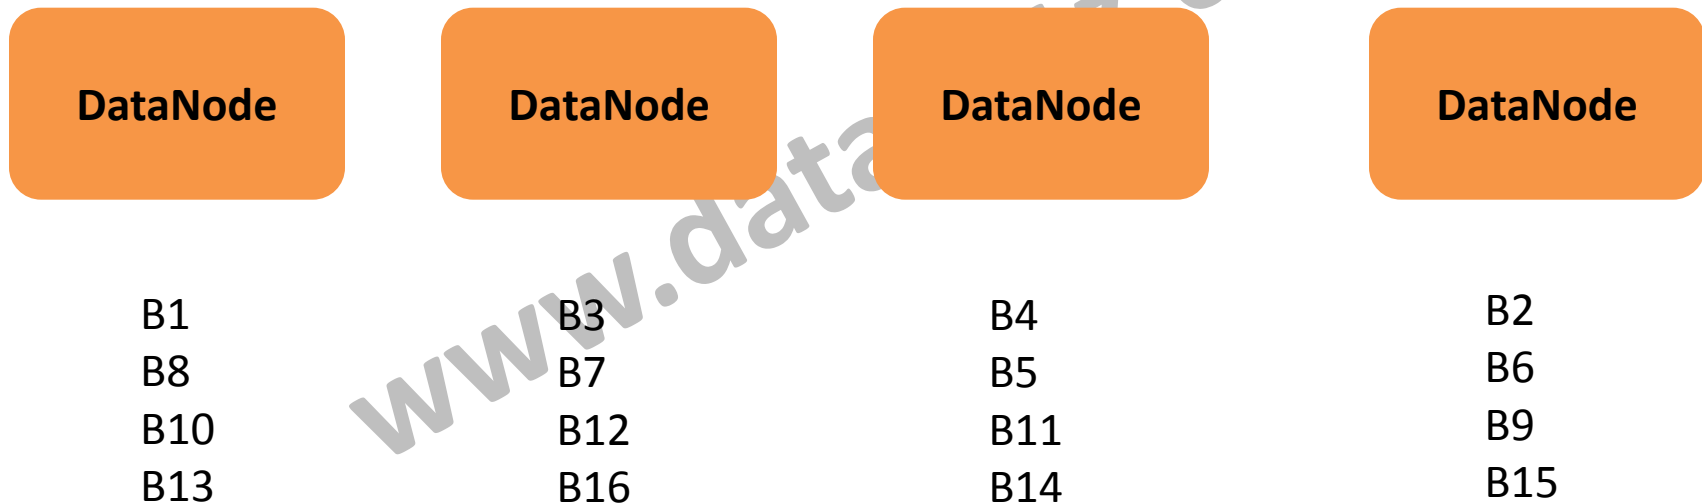
\* - will be seen later

# Block Concept

TestFile1.txt    ->    1GB  
Block Size        ->    64 MB

Files are splitted into number of  
chunks(Blocks) of pre-defined size

No of Blocks =  $1\text{GB} / 64\text{MB} = 16$  blocks  
Blocks are B1,B2,.....B16

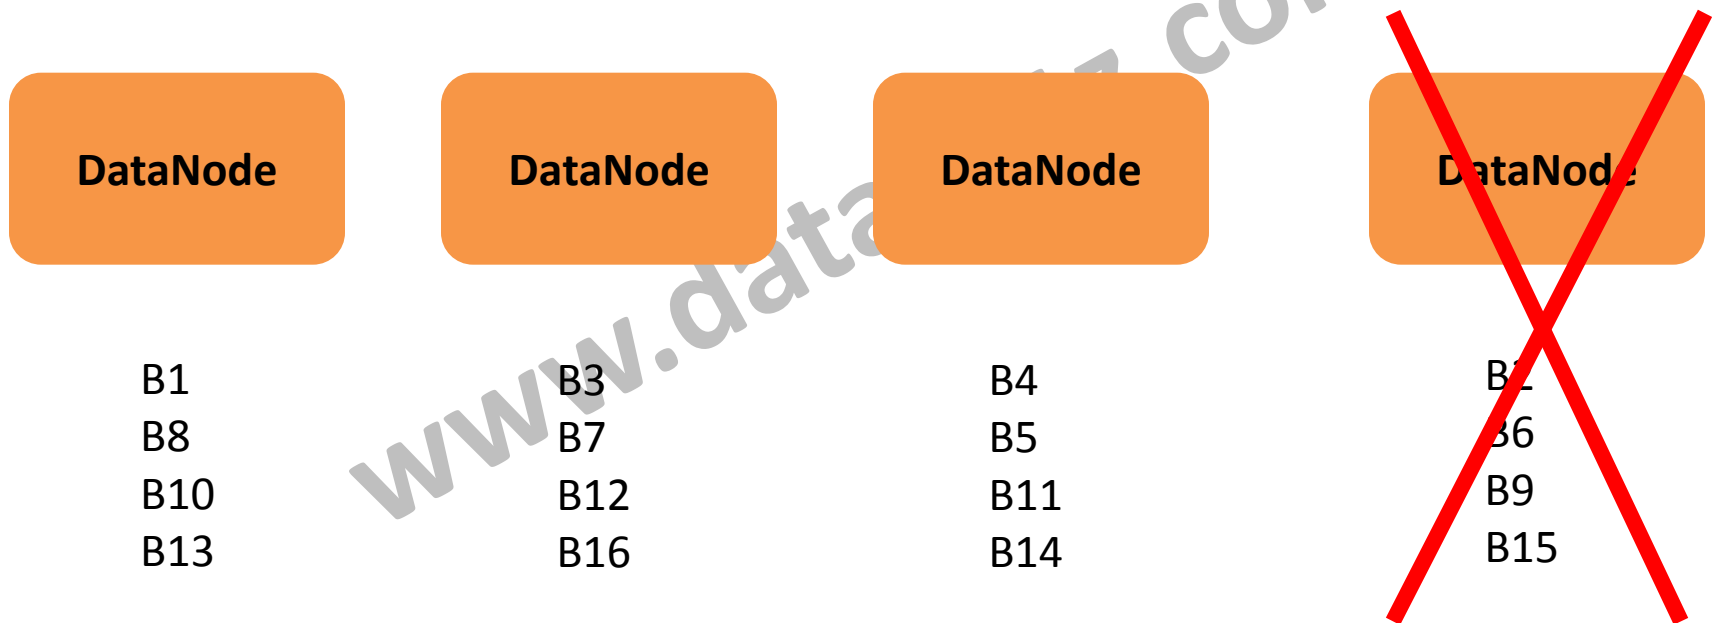


# Block Concept

TestFile1.txt -> 1GB  
Block Size -> 64 MB

No of Blocks =  $1\text{GB} / 64\text{MB} = 16$  blocks  
Blocks are B1,B2,.....B16

**What happens to my data if  
node 4 goes down??**



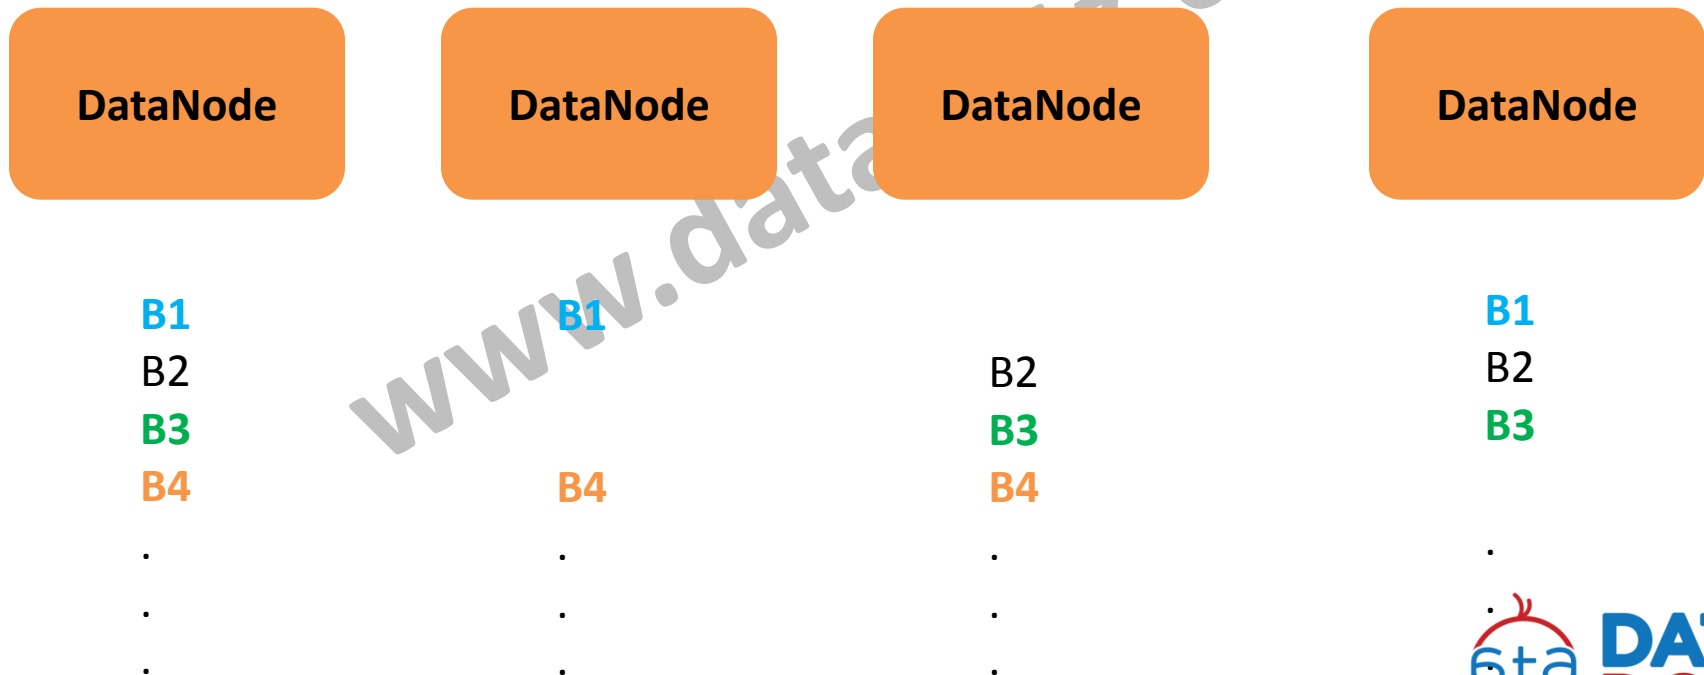


# Fault Tolerant in HDFS

TestFile1.txt -> 1GB  
Block Size -> 64 MB

No of Blocks =  $1\text{GB} / 64\text{MB} = 16$  blocks  
Blocks are B1,B2,.....B16

HDFS provides fault tolerant by replication of each block by 3

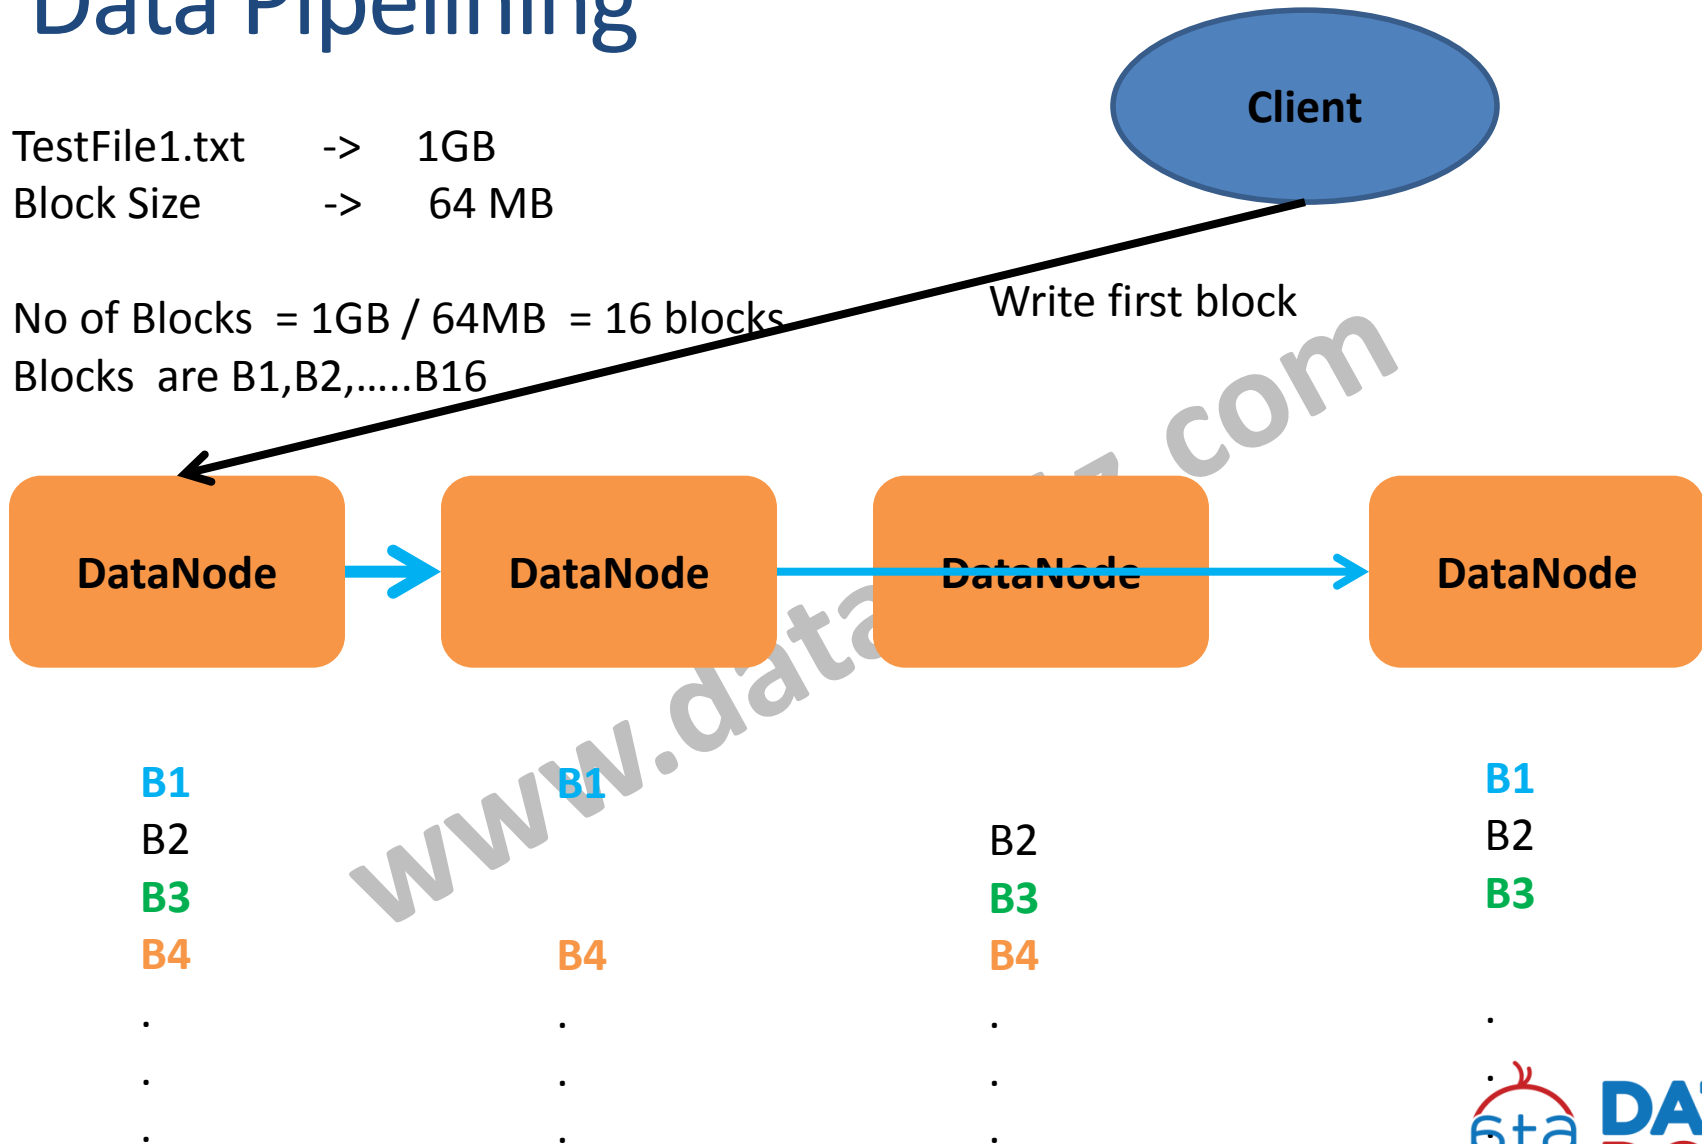


# Data Pipelining

TestFile1.txt -> 1GB  
Block Size -> 64 MB

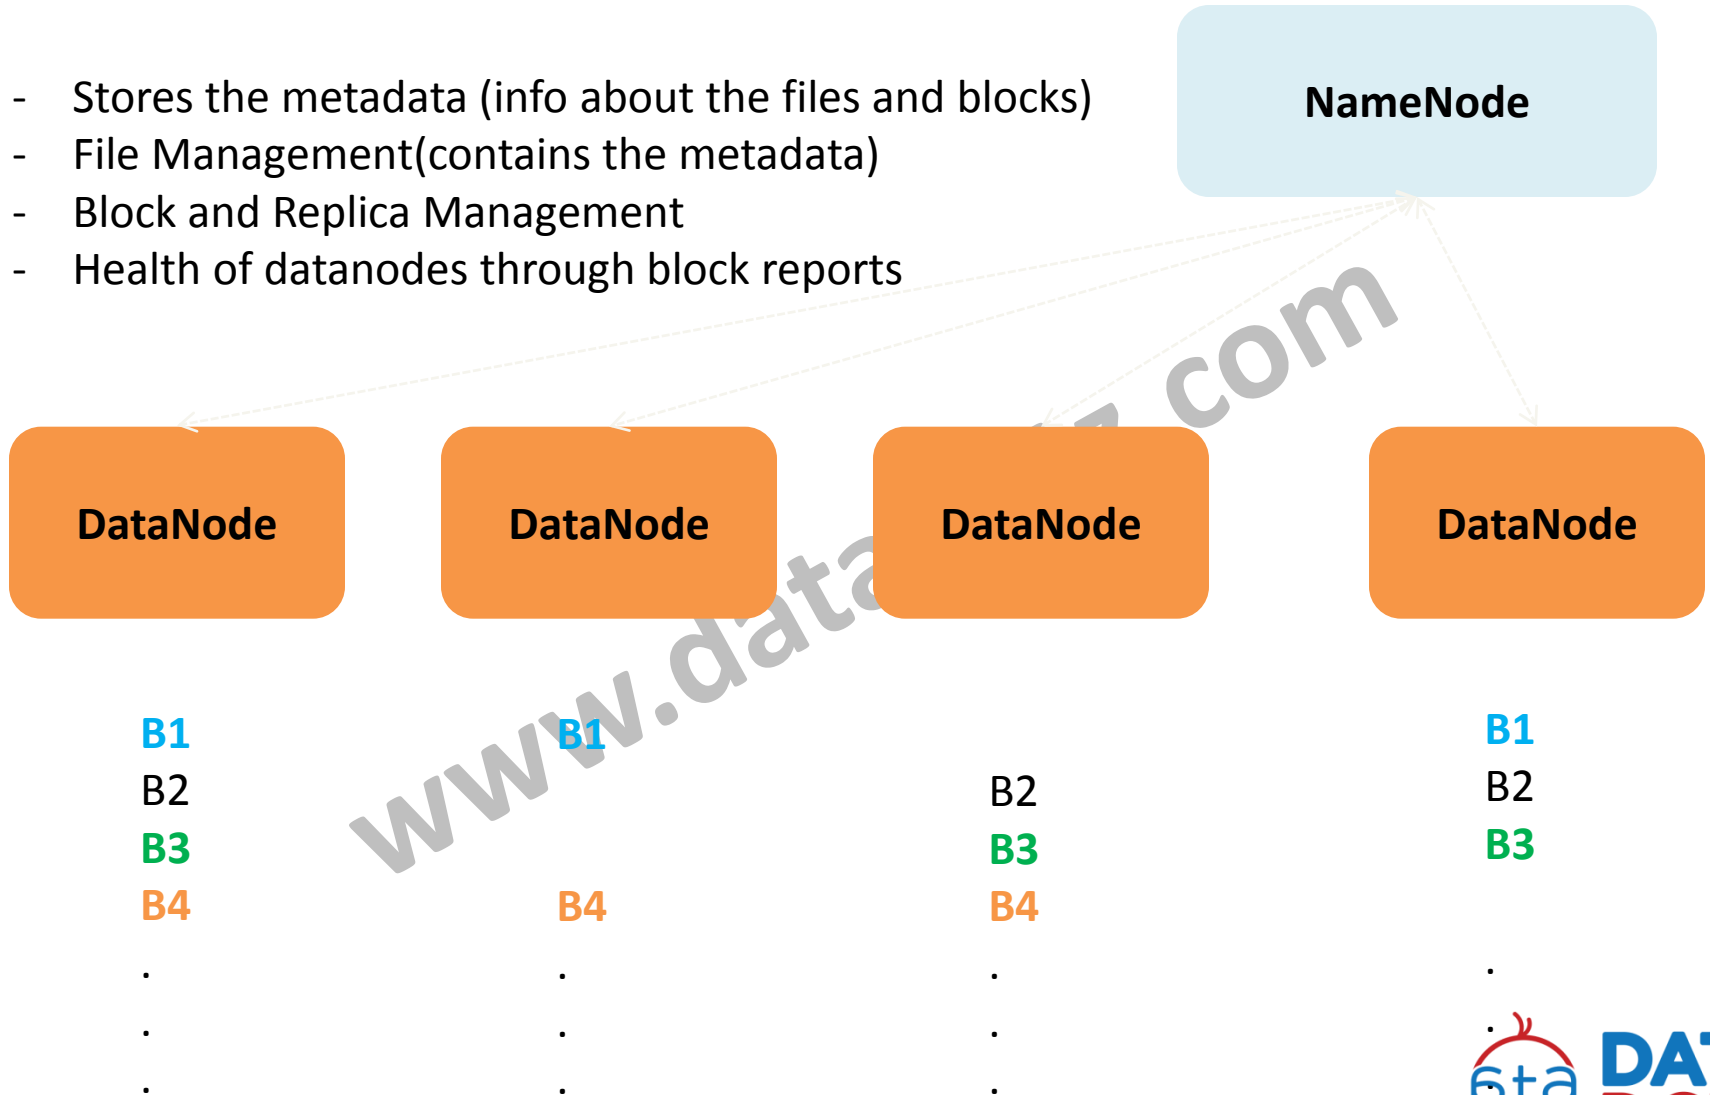
No of Blocks =  $1\text{GB} / 64\text{MB} = 16$  blocks  
Blocks are B1,B2,.....B16

Write first block



# Role of NameNode

- Stores the metadata (info about the files and blocks)
- File Management(contains the metadata)
- Block and Replica Management
- Health of datanodes through block reports



# Execution Modes & Installation

- Modes
  - Single Stand Alone
    - All Process runs in a single jvm
    - Does not use HDFS
  - Pseudo Distributed Mode - **for our training**
    - All daemon process runs in separate jvm in a single local machine
    - Used for development and testing
    - Uses HDFS to store data
  - Distributed Mode
    - A cluster of nodes more than 1
    - Each Process may run in different nodes
- Please follow instructor and doc provided

# Installation

- Please follow the steps in the document given

[www.datadotz.com](http://www.datadotz.com)

# Ports used by Hadoop Daemons

Daemon	RPC	WEB
NameNode	8020 (50000*)	50070
SecondaryNameNode		50090
JobTracker	8021(50001*)	50030
TaskTracker	50020	50060
DataNode	50010	50075

www.data

# After installation

- jps
  - JPS - JVM profiling status tool
- Web UI
  - NameNode - <http://localhost:50070>
  - JobTracker – <http://localhost:50030>

www.datadotz.com

# Accessing HDFS

- Command line
  - Usage: *hadoop dfs <command>*
- JAVA API
- webHDFS

www.datadotz.com



# HDFS commands

- *hadoop dfs -copyFromLocal <srcLOCALfile> <destHDFSfile>*
- *hadoop dfs -ls /*
- *hadoop dfs -cat /<destHDFSfile>*
- *hadoop dfs -copyToLocal <srcHDFSfile> <destLOCALfile>*
- *hadoop dfs -mkdir /test*
- *hadoop dfs -rmr /test*
- *Please follow the document given*

# JAVA API

- Most Packages Used

- *org.apache.hadoop.conf.Configuration*
- *org.apache.hadoop.fs.BlockLocation*
- *org.apache.hadoop.fs.FSDataInputStream*
- *org.apache.hadoop.fs.FSDataOutputStream*
- *org.apache.hadoop.fs.FileStatus*
- *org.apache.hadoop.fs.FileSystem*
- *org.apache.hadoop.fs.Path*
- *org.apache.hadoop.hdfs.DistributedFileSystem*
- *org.apache.hadoop.hdfs.protocol.DatanodeInfo*

**Please see and execute the example code provided**

# FileSystem API methods

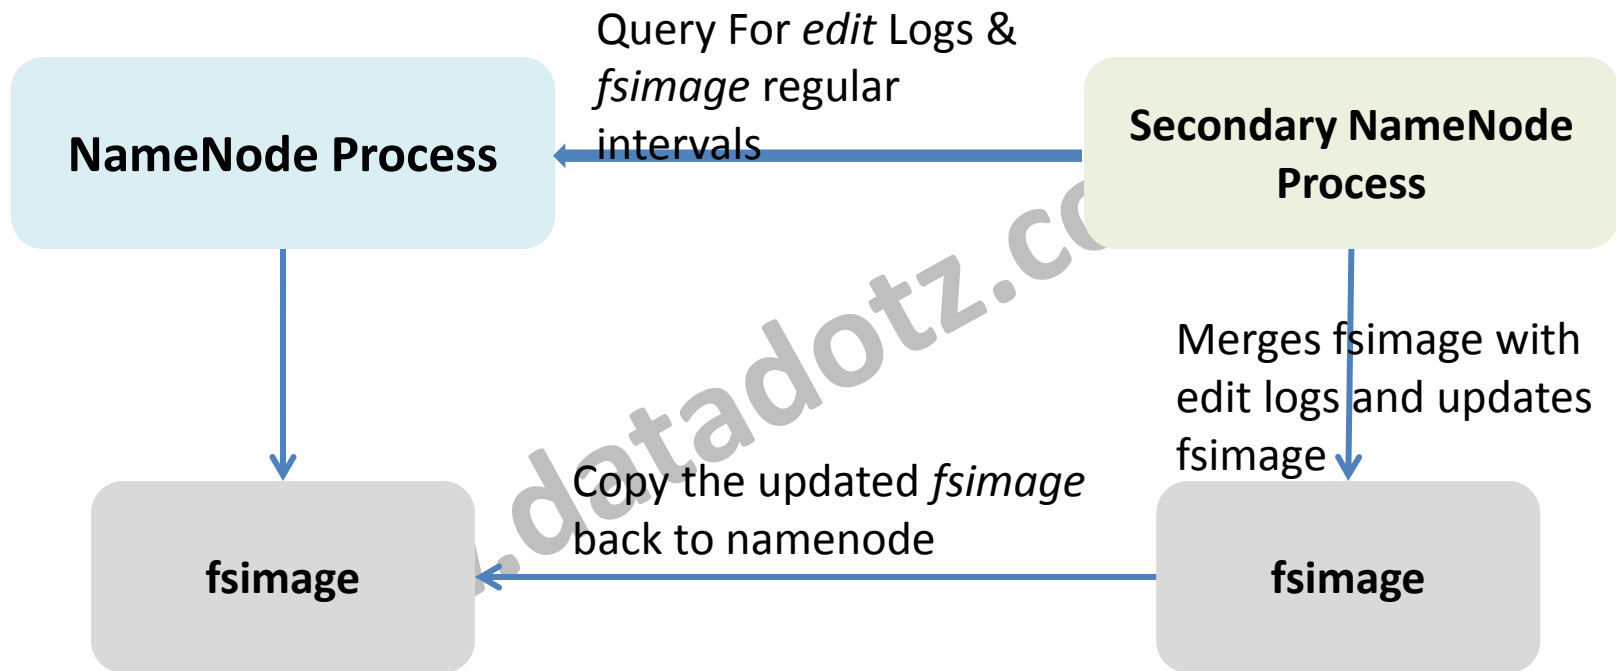
- `append()`
- `copyFromLocalFile()`
- `create()`
- `delete()`
- `makedirs()`
- `open()`

[www.datadotz.com](http://www.datadotz.com)

# Secondary NameNode\*

- A helper node for NameNode
- performs memory-intensive administrative functions for the
- NameNode
- Have a checkpoint for the file system (HDFS)
- Not a Backup Node

# Role of Secondary NameNode



Thank you

